

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: from tensorflow.keras.datasets import cifar10
```

```
In [4]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

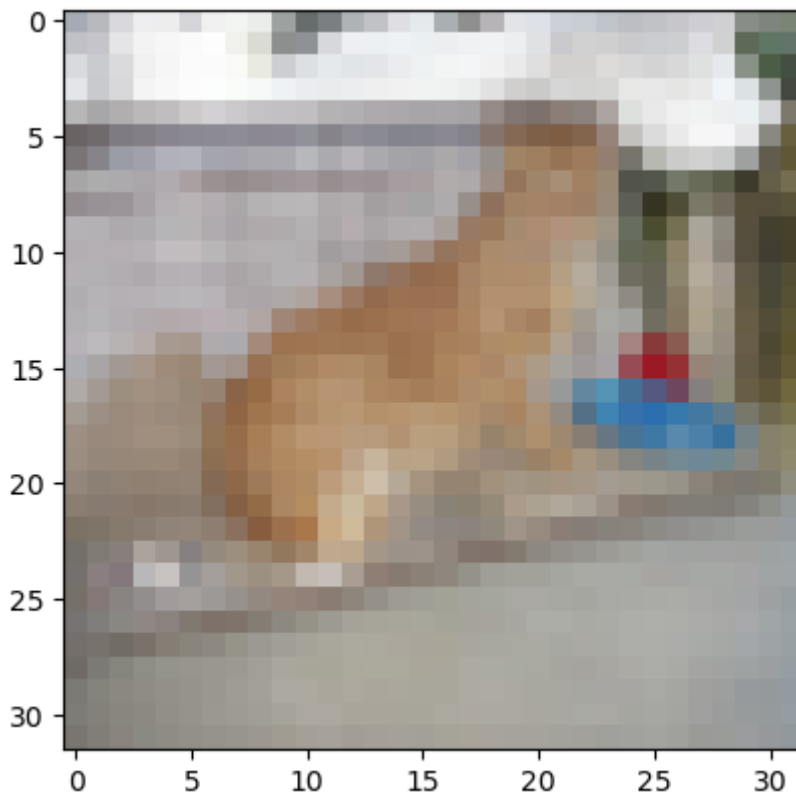
Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 119s 1us/step

```
In [5]: x_train.shape
```

```
Out[5]: (50000, 32, 32, 3)
```

```
In [12]: plt.imshow(x_train[1030])
```

```
Out[12]: <matplotlib.image.AxesImage at 0x1e42efb1be0>
```



```
In [13]: x_train[0].max()
```

```
Out[13]: 255
```

```
In [14]: x_train = x_train/255
```

```
In [15]: x_test = x_test/255
```

```
In [17]: from tensorflow.keras.utils import to_categorical
```

```
In [18]: y_cat_train = to_categorical(y_train, 10)
```

```
In [19]: y_cat_test = to_categorical(y_test, 10)
```

```
In [20]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Dense, MaxPool2D, Flatten
```

```
In [22]: model = Sequential()

model.add(Conv2D(filters=32, kernel_size=(4,4), input_shape=(32,32,3), activation=
model.add(MaxPool2D(pool_size=(2,2)))

model.add(Conv2D(filters=32, kernel_size=(4,4), input_shape=(32,32,3), activation=
model.add(MaxPool2D(pool_size=(2,2)))

model.add(Flatten())

model.add(Dense(256, activation='relu'))

model.add(Dense(10,activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [23]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 29, 29, 32)	1568
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 32)	16416
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 32)	0
flatten (Flatten)	(None, 800)	0
dense (Dense)	(None, 256)	205056
dense_1 (Dense)	(None, 10)	2570
=====		
Total params: 225,610		
Trainable params: 225,610		
Non-trainable params: 0		

```
In [24]: from tensorflow.keras.callbacks import EarlyStopping
```

```
In [25]: early_stop = EarlyStopping(monitor='val_loss', patience=2)
```

```
In [26]: model.fit(x_train, y_cat_train, epochs=15, validation_data=(x_test, y_cat_test), callbacks=[early_stop])
```

```


Epoch 1/15
1563/1563 [=====] - 17s 11ms/step - loss: 1.4979 - accuracy: 0.4560 - val_loss: 1.3093 - val_accuracy: 0.5371
Epoch 2/15
1563/1563 [=====] - 16s 10ms/step - loss: 1.1721 - accuracy: 0.5858 - val_loss: 1.1244 - val_accuracy: 0.6054
Epoch 3/15
1563/1563 [=====] - 16s 10ms/step - loss: 1.0305 - accuracy: 0.6373 - val_loss: 1.0201 - val_accuracy: 0.6520
Epoch 4/15
1563/1563 [=====] - 15s 10ms/step - loss: 0.9310 - accuracy: 0.6759 - val_loss: 1.0059 - val_accuracy: 0.6509
Epoch 5/15
1563/1563 [=====] - 15s 10ms/step - loss: 0.8487 - accuracy: 0.7042 - val_loss: 1.0048 - val_accuracy: 0.6568
Epoch 6/15
1563/1563 [=====] - 15s 10ms/step - loss: 0.7730 - accuracy: 0.7319 - val_loss: 0.9806 - val_accuracy: 0.6714
Epoch 7/15
1563/1563 [=====] - 15s 9ms/step - loss: 0.7070 - accuracy: 0.7531 - val_loss: 0.9803 - val_accuracy: 0.6725
Epoch 8/15
1563/1563 [=====] - 15s 9ms/step - loss: 0.6504 - accuracy: 0.7725 - val_loss: 1.0212 - val_accuracy: 0.6739
Epoch 9/15
1563/1563 [=====] - 15s 9ms/step - loss: 0.5929 - accuracy: 0.7916 - val_loss: 1.0579 - val_accuracy: 0.6663
Out[26]: <keras.callbacks.History at 0x1e423867610>

```

```
In [27]: model_loss = pd.DataFrame(model.history.history)
```


```
In [30]: model_loss[['val_loss', 'loss']].plot()
```

```
Out[30]: <AxesSubplot:>
```



```
In [31]: model_loss[['val_accuracy', 'accuracy']].plot()
```

Out[31]: <AxesSubplot:>



In [32]: `model.evaluate(x_test, y_cat_test, verbose=0)`

Out[32]: [1.0579068660736084, 0.6662999987602234]

In [33]: `from sklearn.metrics import classification_report, confusion_matrix`

In [34]: `predictions = model.predict(x_test)`

313/313 [=====] - 1s 3ms/step

In [37]: `prediction = np.argmax(predictions,axis=1)`

In [38]: `print(classification_report(y_test, prediction))`


	precision	recall	f1-score	support
0	0.69	0.75	0.72	1000
1	0.75	0.83	0.79	1000
2	0.53	0.67	0.59	1000
3	0.49	0.46	0.48	1000
4	0.72	0.49	0.58	1000
5	0.56	0.56	0.56	1000
6	0.84	0.67	0.74	1000
7	0.69	0.70	0.70	1000
8	0.74	0.79	0.77	1000
9	0.71	0.74	0.72	1000
accuracy			0.67	10000
macro avg	0.67	0.67	0.66	10000
weighted avg	0.67	0.67	0.66	10000

In [39]: `confusion_matrix(y_test, prediction)`

```
Out[39]: array([[747, 31, 54, 19, 7, 8, 3, 5, 73, 53],
 [ 18, 833, 9, 7, 2, 12, 5, 5, 39, 70],
 [ 65, 7, 668, 53, 33, 60, 28, 42, 24, 20],
 [ 29, 13, 100, 462, 48, 196, 40, 52, 27, 33],
 [ 39, 15, 168, 79, 486, 49, 30, 94, 25, 15],
 [ 16, 9, 93, 166, 28, 565, 17, 70, 15, 21],
 [ 14, 26, 91, 75, 23, 45, 670, 27, 14, 15],
 [ 35, 5, 59, 38, 39, 69, 3, 703, 9, 40],
 [ 82, 49, 14, 16, 3, 7, 0, 3, 790, 36],
 [ 36, 120, 13, 19, 5, 5, 5, 13, 45, 739]], dtype=int64)
```

```
In [40]: sns.heatmap(confusion_matrix(y_test, prediction))
```


```
Out[40]: <AxesSubplot:>
```



```
In [46]: my_image = x_test[34]
```

```
In [47]: plt.imshow(my_image)
```

```
Out[47]: <matplotlib.image.AxesImage at 0x1e425b59dc0>
```



```
In [48]: np.argmax(model.predict(my_image.reshape(1,32,32,3)),axis=1)
```

```
1/1 [=====] - 0s 14ms/step
```

```
Out[48]: array([9], dtype=int64)
```

```
In [ ]:
```