

Deep Learning for Image Classification Assessment

Building an image classifier with Keras and Convolutional Neural Networks for the Fashion MNIST dataset. This data set includes 10 labels of different clothing types with 28 by 28 *grayscale* images. There is a training set of 60,000 images and 10,000 test images.**

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

The Data


```
In [56]: from tensorflow.keras.datasets import fashion_mnist  
  
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

Visualizing the Data

```
In [57]: import matplotlib.pyplot as plt  
%matplotlib inline
```

```
In [116... plt.imshow(x_train[0])
```

```
Out[116]: <matplotlib.image.AxesImage at 0x2a1da9b2910>
```



Preprocessing the Data

```
In [59]: x_train.max()  
x_test.max()
```

```
Out[59]: 255
```

```
In [60]: x_train = x_train/255
```

```
In [61]: x_test = x_test/255
```

```
In [62]: x_train.max()
```

```
Out[62]: 1.0
```

```
In [63]: import numpy as np
```

```
In [64]: x_train = x_train.reshape(60000,28,28,1)  
x_test = x_test.reshape(10000,28,28,1)
```

```
In [65]: x_train.shape
```

```
Out[65]: (60000, 28, 28, 1)
```

```
In [66]: x_test.shape
```

```
Out[66]: (10000, 28, 28, 1)
```

```
In [67]: from tensorflow.keras.utils import to_categorical
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js)

```
In [69]: y_cat_test = to_categorical(y_test)
```

```
In [70]: y_cat_train
```

```
Out[70]: array([[0., 0., 0., ..., 0., 0., 1.],
               [1., 0., 0., ..., 0., 0., 0.],
               [1., 0., 0., ..., 0., 0., 0.],
               ...,
               [0., 0., 0., ..., 0., 0., 0.],
               [1., 0., 0., ..., 0., 0., 0.],
               [0., 0., 0., ..., 0., 0., 0.]], dtype=float32)
```

Building the Model

```
In [71]: from tensorflow.keras.models import Sequential
         from tensorflow.keras.layers import Conv2D, Flatten, MaxPool2D, Dense
```

```
In [108... model = Sequential()

model.add(Conv2D(filters=32, kernel_size=(4,4),input_shape=(28, 28, 1), activation='relu'))
model.add(MaxPool2D(pool_size=(2,2)))

model.add(Flatten())

model.add(Dense(128, activation='relu'))

model.add(Dense(10, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

```
In [109... model.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
conv2d_14 (Conv2D)	(None, 25, 25, 32)	544
max_pooling2d_14 (MaxPooling2D)	(None, 12, 12, 32)	0
flatten_7 (Flatten)	(None, 4608)	0
dense_14 (Dense)	(None, 128)	589952
dense_15 (Dense)	(None, 10)	1290
=====		
Total params: 591,786		
Trainable params: 591,786		
Non-trainable params: 0		

Training the Model

```
In [110... model.fit(x_train,y_cat_train, epochs=10)
```

```

Epoch 1/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.4193 - accurac
y: 0.8497
Epoch 2/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.2818 - accurac
y: 0.8981
Epoch 3/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.2402 - accurac
y: 0.9115
Epoch 4/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.2120 - accurac
y: 0.9230
Epoch 5/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1897 - accurac
y: 0.9313
Epoch 6/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1725 - accurac
y: 0.9372
Epoch 7/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1560 - accurac
y: 0.9449
Epoch 8/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1445 - accurac
y: 0.9479
Epoch 9/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1319 - accurac
y: 0.9535
Epoch 10/10
1875/1875 [=====] - 11s 6ms/step - loss: 0.1203 - accurac
y: 0.9569
Out[110]: <keras.callbacks.History at 0x2a1f176adf0>

```

Evaluating the Model

```
In [111... from sklearn.metrics import classification_report, confusion_matrix
```

```
In [112... predictions = model.predict(x_test)
313/313 [=====] - 1s 2ms/step
```

```
In [113... predictions_arg = np.argmax(predictions, axis=1)
```

```
In [114... predictions_arg[0]
```

```
Out[114]: 9
```

```
In [115... print(classification_report(y_test, predictions_arg))
```

	precision	recall	f1-score	support
0	0.83	0.91	0.87	1000
1	0.98	0.98	0.98	1000
2	0.83	0.89	0.86	1000
3	0.89	0.94	0.91	1000
4	0.89	0.81	0.85	1000
5	0.97	0.98	0.98	1000
6	0.80	0.70	0.75	1000
7	0.94	0.98	0.96	1000
8	0.98	0.98	0.98	1000
9	0.98	0.95	0.96	1000
accuracy			0.91	10000
macro avg	0.91	0.91	0.91	10000
weighted avg	0.91	0.91	0.91	10000